

# Out-of-Core Property Graph Matching for Real-World Graph Databases

Yanxuan Cui  
Tsinghua University  
Beijing, China  
cuiyx18@mails.tsinghua.edu.cn

## ABSTRACT

### PVLDB Reference Format:

Yanxuan Cui. Out-of-Core Property Graph Matching for Real-World Graph Databases. PVLDB, 14(1): XXX-XXX, 2021.  
doi:XX.XX/XXX.XX

### PVLDB Artifact Availability:

The source code, data, and/or other artifacts have been made available at [http://vldb.org/pvldb/format\\_vol14.html](http://vldb.org/pvldb/format_vol14.html).

## 1 INTRODUCTION

## 2 BACKGROUND

This section introduces the formal definition of property graphs, and then discusses the property graph matching problem.

### 2.1 Property Graph Model

A *property graph* is a directed vertex-labeled edge-labeled multi-graph with self-edges, and key-value properties are stored on vertices and edges. We now provide the formal definition of a property graph.

**DEFINITION 1 (PROPERTY GRAPH [1]).** A property graph  $G$  is a tuple  $(V, E, \rho, \lambda, \sigma)$ , where:

- (1)  $V$  is a finite set of vertices.
- (2)  $E$  is a finite set of edges such that  $V$  and  $E$  have no elements in common.
- (3)  $\rho : E \rightarrow (V \times V)$  is a total function. Intuitively,  $\rho(e) = (v_1, v_2)$  indicates that  $e$  is a directed edge from  $v_1$  to  $v_2$ .
- (4)  $\lambda : (V \cup E) \rightarrow L$  is a total function where  $L$  is a set of labels. Intuitively, if  $v \in V$ ,  $\rho(v) = l$  (respectively,  $e \in E$ ,  $\rho(e) = l$ ), then  $l$  is the label of vertex  $v$  (respectively, edge  $e$ ).
- (5)  $\sigma : (V \cup E) \times Prop \rightarrow Val$  is a partial function with  $Prop$  a finite set of properties and  $Val$  a set of values. Intuitively, if  $v \in V$ ,  $p \in Prop$ ,  $\sigma(v, p) = s$  (respectively,  $e \in E$ ,  $p \in Prop$ ,  $\sigma(e, p) = s$ ), then  $s$  is the value of property  $p$  for vertex  $v$  (respectively, edge  $e$ ) in the property graph  $G$ .

For simplicity, in this paper, we do not discuss the properties i.e.,  $\sigma$  in  $G$ , because similar techniques can be used as processing the labels  $\lambda$ . Thus, the property graph  $G$  can be denoted by

$(V(G), E(G), \rho_G, \lambda_G)$ . Please note that the total function  $\rho_G$  is necessary, in general, we cannot identify an edge simply by the starting and ending vertices such as  $(u_1, u_2)$  as can be done in the simple graph model, because multiple edges may appear between the two vertices. However, we may use the  $(u_1, u_2)$  notation if all we care about is that there exist at least one edge between  $u_1$  and  $u_2$ .

**DEFINITION 2 (VERTEX COVER).** A vertex cover  $V_c$  of a property graph  $G$  is a subset of  $V(G)$  such that  $\forall e \in E(G), \rho_G(e) = (u, v) \implies u \in V_c \vee v \in V_c$ .

### 2.2 Property Graph Matching Problem

**DEFINITION 3 (SUBGRAPH).** A property graph  $F$  is called a subgraph of a property graph  $G$ , written  $F \subseteq G$ , if  $V(F) \subseteq V(G)$ ,  $E(F) \subseteq E(G)$ ,  $\rho_F$  is a restriction of  $\rho_G$ , and  $\lambda_F$  is a restriction of  $\lambda_G$ .

Let  $G$  be any property graph, and let  $S \subseteq V(G)$ , then the *induced subgraph*  $G[S]$  is the graph whose vertex set is  $S$  and whose edge set consists of all of the edges in  $E(G)$  that have both endpoints in  $S$ .

**DEFINITION 4 (PROPERTY GRAPH ISOMORPHISM).** Two property graphs  $G$  and  $H$  are isomorphic, written  $G \cong H$ , if there exists bijections  $\theta : V(G) \rightarrow V(H)$  and  $\phi : E(G) \rightarrow E(H)$  such that  $\rho_G(e) = (u, v)$  if and only if  $\rho_H(\phi(e)) = (\theta(u), \theta(v))$ ,  $\lambda_G(v) = \lambda_H(\theta(v))$  for all  $v \in V(G)$  and  $\lambda_G(e) = \lambda_H(\phi(e))$  for all  $e \in E(G)$ ; Such a pair of mappings is called an isomorphism between  $G$  and  $H$ .

The bijection  $\theta : V(G) \rightarrow V(H)$  is the key in the definition of property graph isomorphism, because the bijection  $\phi : E(G) \rightarrow E(H)$  is straightforward if  $\theta$  is fixed. However, due to automorphism, where an *automorphism* of a graph is an isomorphism of the graph to itself, the bijection  $\theta$  may not be unique.

**DEFINITION 5 (PROPERTY GRAPH MATCHING).** Given a data property graph  $D$ , a pattern property graph  $P$  and a searching condition  $\psi : PG \rightarrow B$  with  $PG$  the set of property graph and  $B$  the set of Boolean values, the property graph matching problem is to report the set  $\mathcal{I} = \{F | F \subseteq D, F \cong P, \psi(F) = \text{true}\}$ .

Authors of previous works usually omit the searching condition  $\psi$  in their definition of graph matching [9, 12, 16, 20]. And they adopt a loosely related technique called *symmetry-breaking* [5], which ensures there is a unique bijection  $\theta : V(P) \rightarrow V(F)$  by providing a partial order on  $V(P)$  after exploiting the automorphism of  $P$ . However, as we have stated before, the WHERE clause is a ubiquitous part of the query language of a graph database. Users of a real-world graph database usually provide their self-defined searching condition  $\psi$  to filter out unnecessary matchings not only symmetry-breaking conditions. Thus, the searching condition we

This work is licensed under the Creative Commons BY-NC-ND 4.0 International License. Visit <https://creativecommons.org/licenses/by-nc-nd/4.0/> to view a copy of this license. For any use beyond those covered by this license, obtain permission by emailing [info@vldb.org](mailto:info@vldb.org). Copyright is held by the owner/author(s). Publication rights licensed to the VLDB Endowment.  
Proceedings of the VLDB Endowment, Vol. 14, No. 1 ISSN 2150–8097.  
doi:XX.XX/XXX.XX

defined here can be viewed as a super set of symmetry-breaking. We add the searching condition in the definition because it is actually a part of the property graph matching problem, and we also found that it can be decomposed and pushed down to lower phase to boost the evaluation of graph matching (Section 3).

A property graph is always directed. However, in some cases such as friendship, there is no need to pay attention to the directions of the edges. In order to support this kind of relationship, a naive approach is to add a duplicate edge in opposite direction for each edge in the data graph. More elegantly, we allow the pattern graph  $P$  to contain undirected edges. Users can simply ignore the direction by providing undirected edges in  $P$  like in industrial graph databases such as Neo4j.

### 3 PROPERTY GRAPH MATCHING FRAMEWORK

#### 4 AN I/O EFFICIENT VERTEX-CENTRIC STORAGE METHOD FOR PROPERTY GRAPH

The random access problem is a well known hard problem for out-of-core systems, especially for the graph related problem, which is notorious for its poor locality [11]. The traditional way to store graphs on disk is the double list method: Stores the in-edges and out-edges separately for each vertex, via the compressed sparse column (CSC) and the compressed sparse row (CSR) format [15]. However, we find that this storage method have limitations to solve the real-world property graph matching problem: *Random disk accesses are unavoidable when in/out-edges are stored separately.*

To solve the random disk access problem, in this section, we propose a novel property graph storage engine from a vertex-centric point of view. Noticeably, only part of the huge billion node data graph would be read when solving a property graph matching problem, with the help of a few easy to implement indices. And all the disk reads are sequential. Moreover, by using the property graph matching engine that will be discussed in the next section, we would be sure that the huge data graph would be read at most once.


##### 4.1 Scan the Data Graph Sequentially

The property graph matching problem requires the complete connection information between vertices, however, the conventional double list storage method break up the completeness and results in random disk accesses.

Consider the patten graph in Figure 1, which simply finds friendship pairs in a data graph. If two vertices in the data graph, say  $v_i$  and  $v_j$ , are supposed to match the pattern, we must be sure that  $v_i$  and  $v_j$  follows each other at the same time. This kind of neighbor checking is also the most essentially building block of a property graph matching engine.

In the data graph,  $v_1$  has  $m$  in-edges and  $n$  out-edges, and suppose that they are stored separately in the traditional way. In order to determine whether  $v_1$  could match  $u_1$ , one has to scan the in-edges (or equally, the out-edges) of  $v_1$  and then check whether the visited neighbor is in the out-edges (or in-edges) list. This scan and check method would significantly slow down the graph matching process,

because the checking process results in random disk read. For real world power-law graphs, where the celebrities or trending topics have a huge amount of followers, the in/out-edges lists stored on disk have to be swapped in and out frequently during the scan and check process as a result of the random disk access pattern. And it would be more complicated if there are more than two edges between  $u_1$  and  $u_2$  in the pattern graph.



**Figure 1: A celebrity with a huge amount of followers and a pattern graph to find friendship pairs in it.**

The source of the problem is that the in/out-edges are stored separately in the conventional method, whereas the property matching process need to retrieve both in- and out- connection information to determine whether a neighbor vertex could be matched. Based on this observation, we propose a vertex-centric property graph storage method that always keeps the necessary connection information together for graph matching. Take Figure 2 as an example, which shows the logical storage structure for the celebrity in Figure 1. Real-world property graphs are directed labeled graphs, instead of splitting the edges into in/out-edges, we treat all the neighbors equally and store the connection information together with the neighbor vertex. The specific edge information (direction, types) could be obtained by the position of the stored edge labels, i.e., the upper labels means the direction is pointed to the root and the lower labels means the opposite. For example, the “:FOLLOWS” at the upper right corner of  $v_3$  means  $v_3$  follows  $v_1$ , the two “:FOLLOWS” besides  $v_{m+1}$  indicate that  $v_1$  and  $v_{m+1}$  follow each other. And multiple edges are ubiquitous in property graphs, we support that by store the edge labels in a sequence, as is shown in the figure where  $v_1$  publishes a social media  $v_{m+n}$  and also likes it.

By storing the neighbors in our vertex-centric approach, all the necessary information are now stored locally together with the vertices, and the neighborhood checking process of a property graph matching engine could be accomplished efficiently within a sequential disk scan. For the pattern in Figure 1, if we suppose that  $v_1$  matches  $u_1$  and want to check whether the neighbors of  $v_1$  would match  $u_2$ , we could just scan the neighbors and check the labels sequentially and find that  $v_{m+1}$  would match. No random disk access appears during the whole process, and there is no need to use complex indices to check the edge labels.

##### 4.2 Simple Indices to Reduce I/Os

Despite of the fact that the size of the whole data graph could be gigantic, we may only care a fraction of the graph by specify the labels in our query for a concrete graph matching problem. There are two basic operations for a graph database to solve a graph

matching problem: 1. Given a data vertex, retrieve the neighbors of it with a specific label; 2. Given a vertex label, retrieve the data vertices with the label. In the following paragraphs, we'll show how to make the two operations efficiently by adding a few simple but efficient indices to the vertex-centric storage engine, which reduces the searching space and I/O significantly.

Consider again the pattern graph in Figure 1, where we only care about the users in the social network, no matter how much social media a user have published or viewed we could just ignore all of them in our specific problem. A straightforward idea is to group the neighbor vertices with the same label together when storing the vertices. However, if the neighbor vertices were stored in the traditional in/out-edge double lists method, we have to group them twice and then still face the information insufficient problem as we discussed in the previous subsection. As for our vertex-centric storage method, the index could be added to the storage engine easily as is shown in Figure 2, where we simply store the key-value pairs that maps the vertex labels to the starting/ending position of the neighbors on disk. Since it is used to locate only the neighbors of a specific vertex, we refer to this kind of index as *local index*. With the help of local index, we could skip all the useless neighbor vertices and only scan the necessary ones.

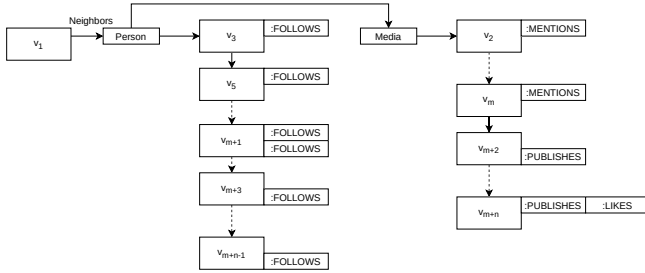


Figure 2: The vertex-centric storage structure of the celebrity  $v_1$  in Figure 1.

From a higher perspective, in order to solve a property graph matching problem, one need to firstly select a starting vertex in the data graph, regardless of the concrete graph matching algorithm whether it is tree based or join based. There is no need to scan all the vertices in a billion node data graph if we only care about vertices with the specific labels. Just like the local index we discussed above, we could add a global index which contains key-value pairs that maps the labels to the corresponding position on disk (Figure 3).

In summary, for a property graph matching problem, we could quickly jump to the domain of interest with the help of the global index, and then scan and check only the necessary neighbors with the help of our local index. After jump to the correct position, all the disk reads are sequential.

### 4.3 Remarks on the Implementation of the Storage Method

For applications that need to perform the two basic operations (visiting vertices and visiting the neighbors), we define two iterators as the interface to visit the data graph:




Figure 3: The storage structure overview of the data graph in Figure 1.

- (1) VERTEXITER: Given a vertex label, it iterates through the vertices with the specified label ordered by the ID of vertices;
- (2) NEIGHBORITER: Given a data vertex and a label, it visits the neighbors with the label, the neighbors are also sorted by the IDs.

In Section 5 we'll show that the sorting constraint could boost the matching of property graphs.

These iterators could be implemented efficiently by using our vertex-centric storage model. And we also provide a compact disk format implementing the vertex-centric storage model in Figure 4. Vertices are sorted by their IDs, and we store in/out-degrees as early filters when scanning the vertices. Edge labels are stored as integers here, however, bitmap could also be used for higher performance. The VERTEXITER just searches the global index and then visits the disk data sequentially; The NeighborIter scans the neighbors with the help of the local index.

Please note that the vertex-centric storage model is not restricted to this disk format. In fact, as long as a storage engine could implement the two iterators efficiently, it could implement the vertex-centric storage model well. For example, the sorted vertices could be replaced with B-tree to make the insertion/deletion operation easier for dynamic graphs. It is also possible to implement the vertex-centric storage model in memory as buffer cache for existing graph databases to achieve better locality. Moreover, we are now working on implementing the vertex-centric storage model on top of relational database to embrace the power of the half-century-old mature technology.

## 5 A ECONOMICAL PROPERTY GRAPH MATCHING ENGINE

### 5.1 Using Stars to Sequentially Scan the Huge Graph at Most Once

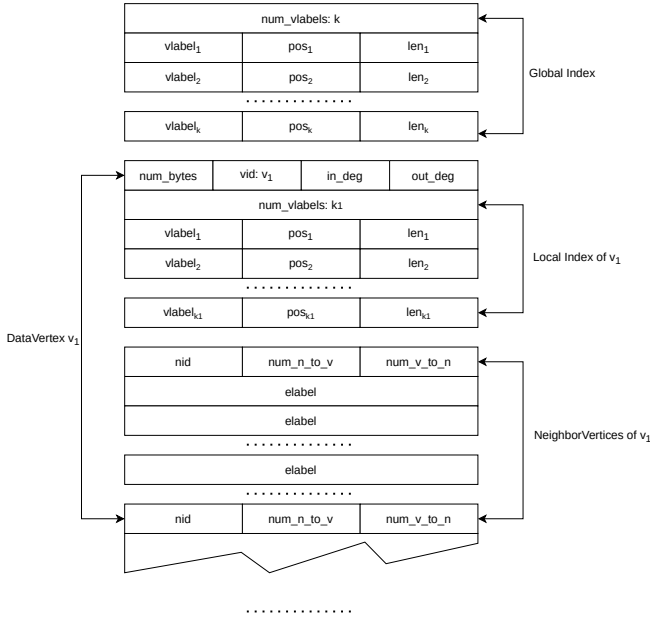
### 5.2 Compressed Matching Results Wrote Sequentially

### 5.3 Scalable Pipeline Join on the Compressed Data

## 6 EXPERIMENTS

## 7 RELATED WORK

Despite the fact that general property graph matching problem is seldom discussed in previous works, simple graph matching has been widely studied. We survey these relevant work in this section.



**Figure 4: An I/O efficient property graph storage method that can boost the graph matching process.**

## In-memory Methods

Most of the early work assumes that the data graph and indices are fit in the main memory of a single machine. Sparked by Ullmann’s backtracking algorithm [22], many subgraph matching algorithms have been proposed using different join order, filter rules, and neighborhood indices [3, 6, 7, 14, 18]. These algorithms usually use a DFS-style tree-based graph exploration to search the matchings without materializing intermediate results. However, these single machine in-memory algorithms are no longer suitable for nowadays billion-node graphs.

To address the scalability problem of single machine in-memory algorithms, many distributed subgraph matching algorithms have been proposed [12, 13, 17, 20, 21]. Because the vertices of the data graph are scattered among machines, these algorithms usually match smaller patterns and get the final result by join operation. For example, Sun et al. [21] introduce a star-like basic matching unit called STwig, and implement their subgraph matching algorithm on top of the Trinity [19] memory cloud. Lai et al. [12] propose TwinTwig join using MapReduce, where a TwinTwig is either a single edge or two incident edges of a vertex. The SEED [13] algorithm use both star and clique as the join units, and use clique compression technique to further improve the performance. However, these distributed algorithms still suffer from severe memory crisis, because the size of partial results grow exponentially with respect to the size of the data graph. Moreover, they must be transferred to other machines before join, which is the most expensive operation in a parallel system such as MapReduce.

Besides, the optimization of a subgraph matching algorithm relies heavily on the underlying graph model:

Unlabeled undirected simple graph is perhaps the simplest graph model, which can be viewed as a special case of property graph

with all the vertices and edges have the same label and have no multi-edges. Some authors distinguish this kind of graphs from others and designate the matching problem of this kind of graph as *subgraph listing* [9, 12, 13, 16, 20, 20, 22]. CBF [16] is the state-of-the-art subgraph listing algorithm, which decompose the pattern graph into a several basic structures called *crystals*, and match these basic units with partial results compressed by the VCBC algorithm. However, it is unable to support general property graph because CBF relies on clique listing to match crystals, which implies the equivalence of vertices in a clique (complete graph) and is not the case of property graph model because of labels and direction of edges.

Another widely studied graph model is vertex-labeled undirected simple graph [4, 6, 17, 18, 21]. TurboISO [6], for example, is turbocharged by the concept of *neighborhood equivalence class* (NEC). It outperforms other competitors by safely avoid the permutation of all possible vertices in the same NEC. A NEC is a set of vertices in the pattern graph, where every vertex has the same label and the same set of neighbors. However, things become more complex and make it not suitable for the property graph model. Because one has to check the labels of vertices, labels of edges, directions of edges in order to test the isomorphism of a property graph, and the real-world multigraphs make life even harder.

## Out-of-core Methods

Many out-of-core triangle enumeration algorithms have been proposed [2, 8, 10, 11]. However, all these algorithms only deal with triangulation, a special case of the graph matching problem. Recently, DUALSIM [9] take a further step and is able to match general unlabeled undirected graphs. To avoid the materialization of intermediate results, it fixes the data vertices by fixing a set of disk pages and then find all matchings in these pages. Apparently, every page of the data graph must be swapped in/out many times in order to get the final result, which lead to severe I/O cost. In contrast, our approach will load the pages sequentially at most once, and we can also use the compressed partial results to boost afterward queries.

## 8 CONCLUSION

## REFERENCES

- [1] Renzo Angles, Marcelo Arenas, Pablo Barceló, Aidan Hogan, Juan L. Reutter, and Domagoj Vrgoc. 2017. Foundations of Modern Query Languages for Graph Databases. *ACM Comput. Surv.* 50, 5 (2017), 68:1–68:40. <https://doi.org/10.1145/3104031>
- [2] Shumo Chu and James Cheng. 2011. Triangle listing in massive networks and its applications. In *Proceedings of the 17th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, San Diego, CA, USA, August 21-24, 2011*, Chid Apté, Joydeep Ghosh, and Padhraic Smyth (Eds.). ACM, 672–680. <https://doi.org/10.1145/2020408.2020513>
- [3] Luigi P. Cordella, Pasquale Foggia, Carlo Sansone, and Mario Vento. 2004. A (Sub)Graph Isomorphism Algorithm for Matching Large Graphs. *IEEE Trans. Pattern Anal. Mach. Intell.* 26, 10 (2004), 1367–1372. <https://doi.org/10.1109/TPAMI.2004.75>
- [4] Vinicius Vitor dos Santos Dias, Carlos H. C. Teixeira, Dorgival O. Guedes, Wagner Meira Jr., and Srinivasan Parthasarathy. 2019. Fractal: A General-Purpose Graph Pattern Mining System. In *Proceedings of the 2019 International Conference on Management of Data, SIGMOD Conference 2019, Amsterdam, The Netherlands, June 30 - July 5, 2019*, Peter A. Boncz, Stefan Manegold, Anastasia Ailamaki, Amol Deshpande, and Tim Kraska (Eds.). ACM, 1357–1374. <https://doi.org/10.1145/3299869.3319875>
- [5] Joshua A. Grochow and Manolis Kellis. 2007. Network Motif Discovery Using Subgraph Enumeration and Symmetry-Breaking. In *Research in Computational Molecular Biology, 11th Annual International Conference, RECOMB 2007, Oakland,*

- CA, USA, April 21-25, 2007, *Proceedings (Lecture Notes in Computer Science)*, Terence P. Speed and Haiyan Huang (Eds.), Vol. 4453, Springer, 92–106. [https://doi.org/10.1007/978-3-540-71681-5\\_7](https://doi.org/10.1007/978-3-540-71681-5_7)
- [6] Wook-Shin Han, Jinsoo Lee, and Jeong-Hoon Lee. 2013. Turbo<sub>iso</sub>: towards ultrafast and robust subgraph isomorphism search in large graph databases. In *Proceedings of the ACM SIGMOD International Conference on Management of Data, SIGMOD 2013, New York, NY, USA, June 22-27, 2013*, Kenneth A. Ross, Divesh Srivastava, and Dimitris Papadias (Eds.). ACM, 337–348. <https://doi.org/10.1145/2463676.2465300>
- [7] Huahai He and Ambuj K. Singh. 2008. Graphs-at-a-time: query language and access methods for graph databases. In *Proceedings of the ACM SIGMOD International Conference on Management of Data, SIGMOD 2008, Vancouver, BC, Canada, June 10-12, 2008*, Jason Tsong-Li Wang (Ed.). ACM, 405–418. <https://doi.org/10.1145/1376616.1376660>
- [8] Xiaocheng Hu, Yufei Tao, and Chin-Wan Chung. 2013. Massive graph triangulation. In *Proceedings of the ACM SIGMOD International Conference on Management of Data, SIGMOD 2013, New York, NY, USA, June 22-27, 2013*, Kenneth A. Ross, Divesh Srivastava, and Dimitris Papadias (Eds.). ACM, 325–336. <https://doi.org/10.1145/2463676.2463704>
- [9] Hyeonji Kim, Juneyoung Lee, Sourav S. Bhowmick, Wook-Shin Han, Jeong-Hoon Lee, Seongyun Ko, and Moath H. A. Jarrah. 2016. DUALSIM: Parallel Subgraph Enumeration in a Massive Graph on a Single Machine. In *Proceedings of the 2016 International Conference on Management of Data, SIGMOD Conference 2016, San Francisco, CA, USA, June 26 - July 01, 2016*, Fatma Özcan, Georgia Koutrika, and Sam Madden (Eds.). ACM, 1231–1245. <https://doi.org/10.1145/2882903.2915209>
- [10] Jinha Kim, Wook-Shin Han, Sangyeon Lee, Kyungyeol Park, and Hwanjo Yu. 2014. OPT: a new framework for overlapped and parallel triangulation in large-scale graphs. In *International Conference on Management of Data, SIGMOD 2014, Snowbird, UT, USA, June 22-27, 2014*, Curtis E. Dyreson, Feifei Li, and M. Tamer Özsu (Eds.). ACM, 637–648. <https://doi.org/10.1145/2588555.2588563>
- [11] Aapo Kyröla, Guy E. Blelloch, and Carlos Guestrin. 2012. GraphChi: Large-Scale Graph Computation on Just a PC. In *10th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2012, Hollywood, CA, USA, October 8-10, 2012*, Chandu Thekkath and Amin Vahdat (Eds.). USENIX Association, 31–46. <https://www.usenix.org/conference/osdi12/technical-sessions/presentation/kyrola>
- [12] Longbin Lai, Lu Qin, Xuemin Lin, and Lijun Chang. 2015. Scalable Subgraph Enumeration in MapReduce. *PVLDB* 8, 10 (2015), 974–985. <https://doi.org/10.14778/2794367.2794368>
- [13] Longbin Lai, Lu Qin, Xuemin Lin, Ying Zhang, and Lijun Chang. 2016. Scalable Distributed Subgraph Enumeration. *PVLDB* 10, 3 (2016), 217–228. <https://doi.org/10.14778/3021924.3021937>
- [14] Jinsoo Lee, Wook-Shin Han, Romans Kasperovics, and Jeong-Hoon Lee. 2012. An In-depth Comparison of Subgraph Isomorphism Algorithms in Graph Databases. *PVLDB* 6, 2 (2012), 133–144. <https://doi.org/10.14778/2535568.2448946>
- [15] Roger A. Pearce, Maya Gokhale, and Nancy M. Amato. 2010. Multithreaded Asynchronous Graph Traversal for In-Memory and Semi-External Memory. In *Conference on High Performance Computing Networking, Storage and Analysis, SC 2010, New Orleans, LA, USA, November 13-19, 2010*. IEEE, 1–11. <https://doi.org/10.1109/SC.2010.34>
- [16] Miao Qiao, Hao Zhang, and Hong Cheng. 2017. Subgraph Matching: on Compression and Computation. *PVLDB* 11, 2 (2017), 176–188. <https://doi.org/10.14778/3149193.3149198>
- [17] Marco Serafini, Gianmarco De Francisci Morales, and Georgos Siganos. 2017. QFrag: distributed graph search via subgraph isomorphism. In *Proceedings of the 2017 Symposium on Cloud Computing, SoCC 2017, Santa Clara, CA, USA, September 24-27, 2017*. ACM, 214–228. <https://doi.org/10.1145/3127479.3131625>
- [18] Haichuan Shang, Ying Zhang, Xuemin Lin, and Jeffrey Xu Yu. 2008. Taming verification hardness: an efficient algorithm for testing subgraph isomorphism. *PVLDB* 1, 1 (2008), 364–375. <https://doi.org/10.14778/1453856.1453899>
- [19] Bin Shao, Haixun Wang, and Yatao Li. 2012. *The Trinity Graph Engine*. Technical Report MSR-TR-2012-30. <https://www.microsoft.com/en-us/research/publication/the-trinity-graph-engine/>
- [20] Yingxia Shao, Bin Cui, Lei Chen, Lin Ma, Junjie Yao, and Ning Xu. 2014. Parallel subgraph listing in a large-scale graph. In *International Conference on Management of Data, SIGMOD 2014, Snowbird, UT, USA, June 22-27, 2014*, Curtis E. Dyreson, Feifei Li, and M. Tamer Özsu (Eds.). ACM, 625–636. <https://doi.org/10.1145/2588555.2588557>
- [21] Zhao Sun, Hongzhi Wang, Haixun Wang, Bin Shao, and Jianzhong Li. 2012. Efficient Subgraph Matching on Billion Node Graphs. *PVLDB* 5, 9 (2012), 788–799. <https://doi.org/10.14778/2311906.2311907>
- [22] Julian R. Ullmann. 1976. An Algorithm for Subgraph Isomorphism. *J. ACM* 23, 1 (1976), 31–42. <https://doi.org/10.1145/321921.321925>